


Data Science Dojo

hadoop



Spark

Big Data Engineering with Distributed Systems

Agenda

- **Introduction:**
 - Data engineering for data scientists
 - The “5 Vs” of Big Data
- **A key problem – machine learning at scale**
- **Distributed computing with Apache Hadoop & Hive**
- **Hadoop in the Azure cloud**
- **Machine learning at scale with Apache Mahout**
- **Distributed computing v2.0 – Apache Spark**

Data Engineering for Data Scientists



Driving a car

VS



Servicing a car

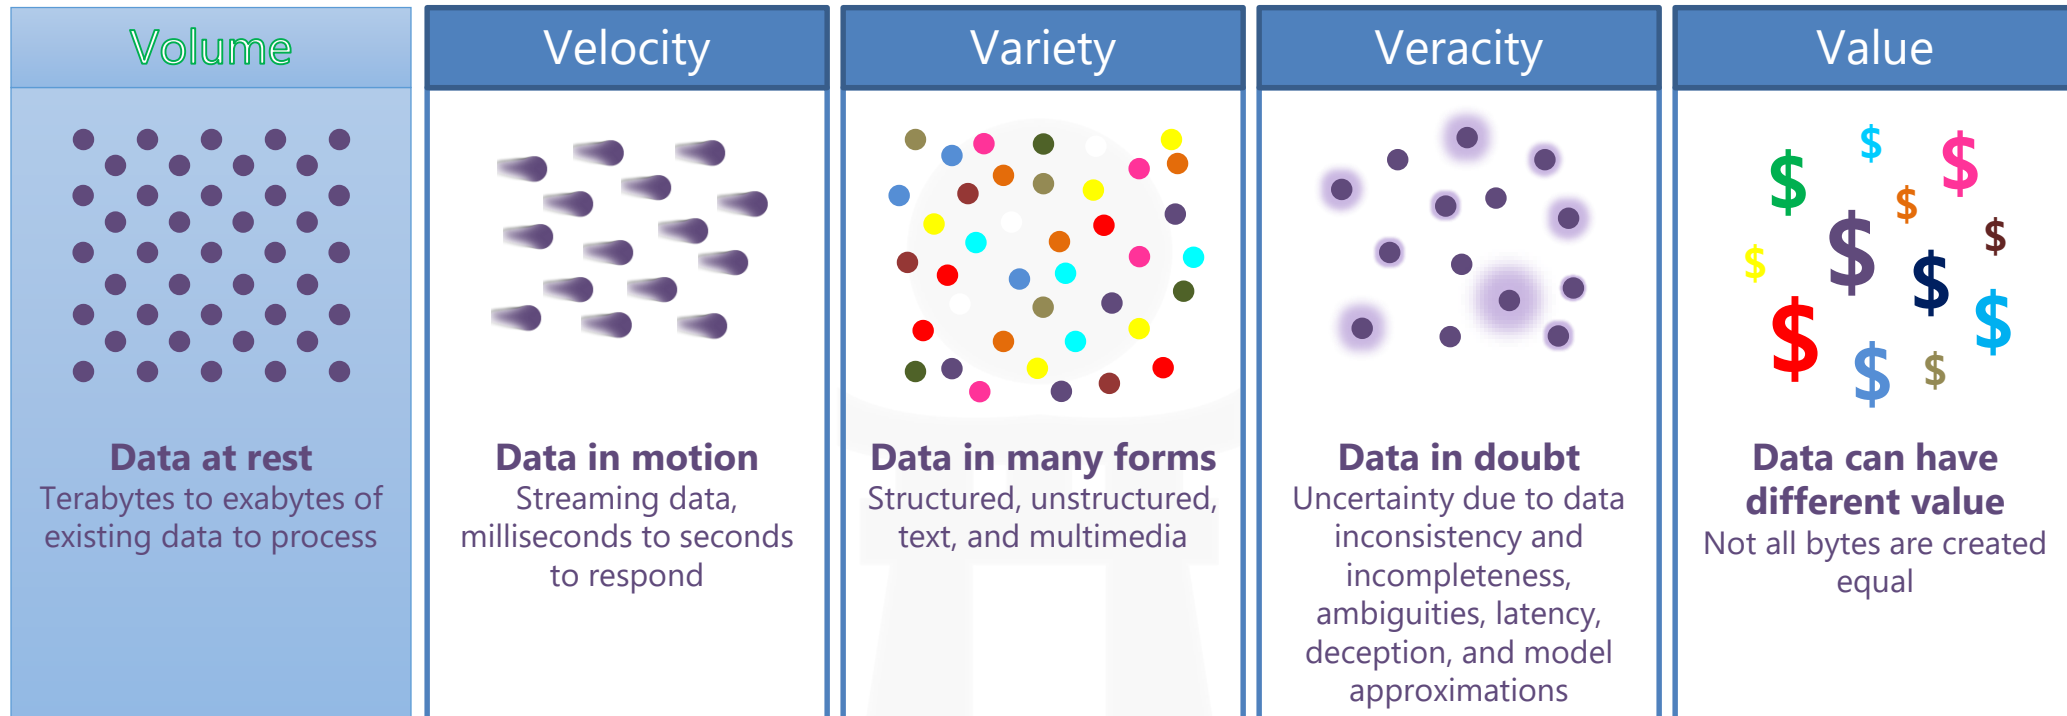
Goals:

- Teach you about data engineering topics/concepts

Non goals:

- Managing or administering a Hadoop cluster

5 Vs of Big Data



- **Goal:** As data scientists we want cost-effective access to the raw materials for our data products!

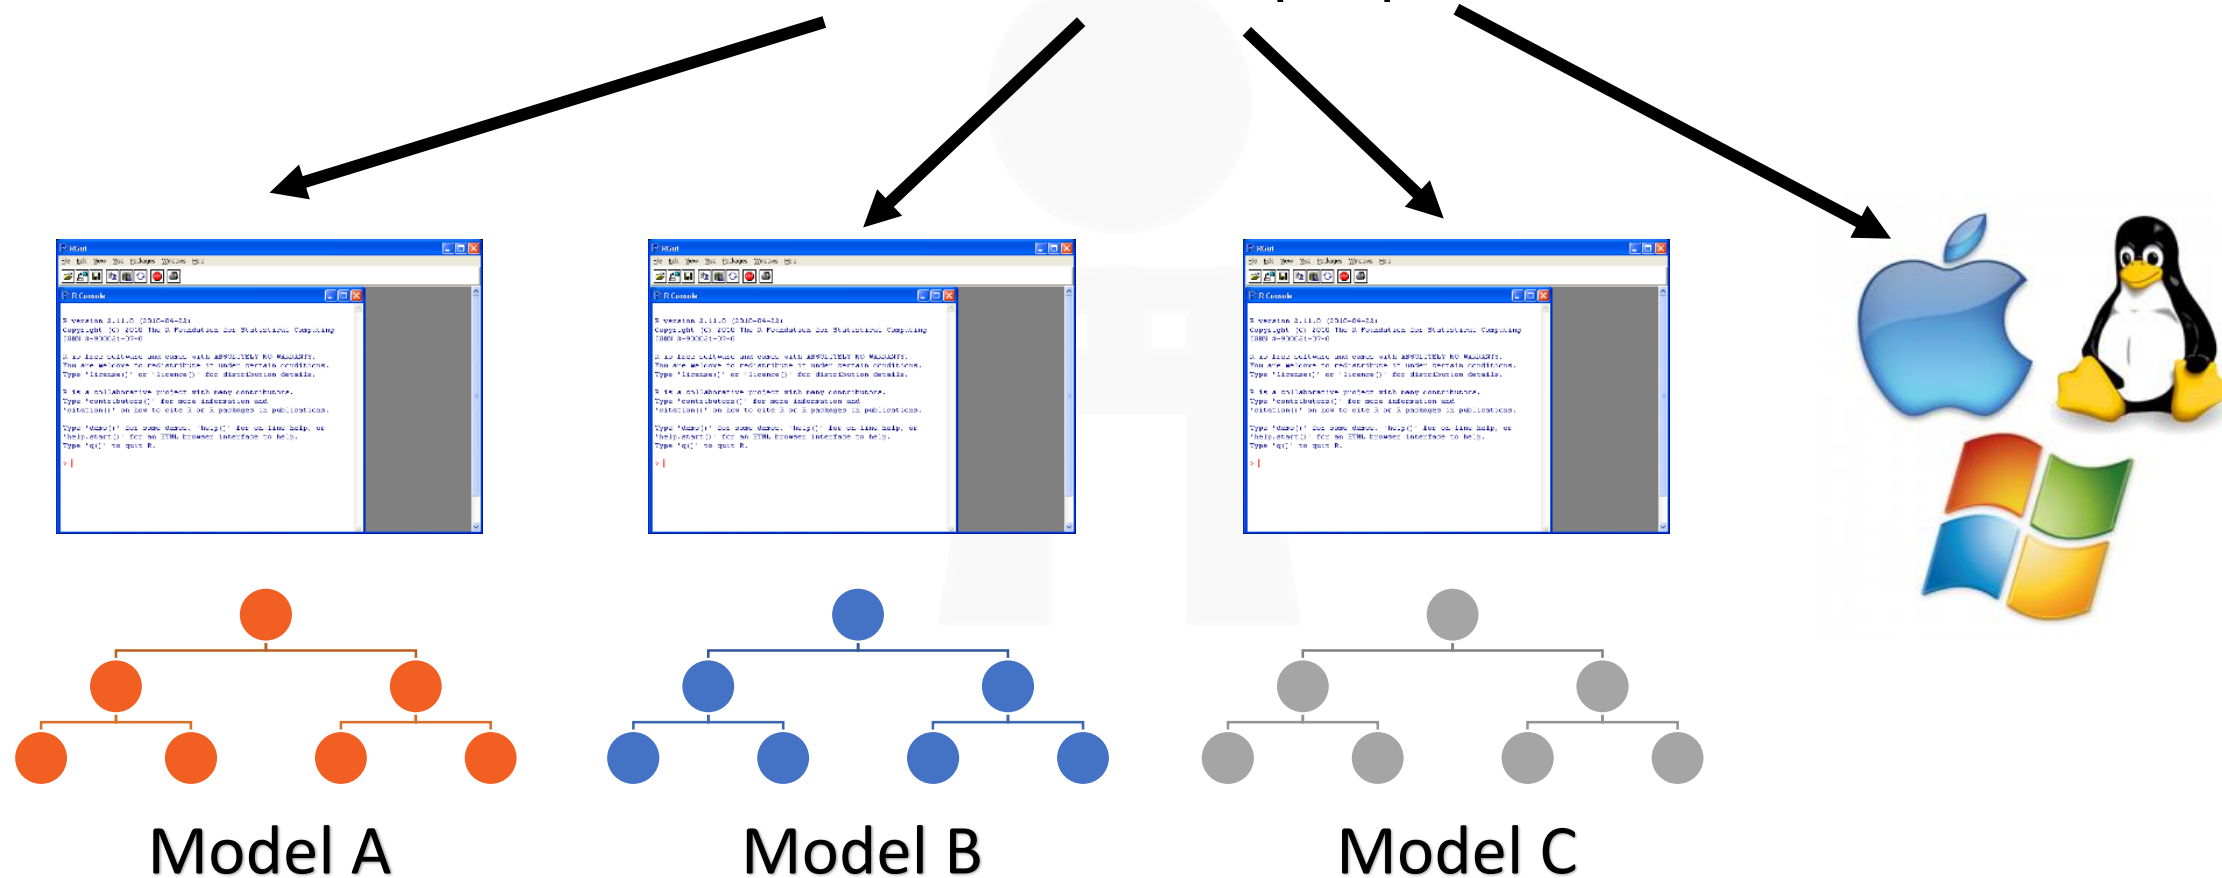


MACHINE LEARNING AT SCALE

OSS R Limits

- Single core
- Single threaded

Quad Core Laptop



OSS R Limits


- Single core
- Single threaded
- All in memory (RAM)
- Vectors & Matrices capped at 4,294,967,295 elements (rows) if 32-bit version; $2^{32} - 1$

OSS R Limits: RAM

- All in memory (RAM)

$$\text{Max Data Limit} = (\text{Total RAM Access} \times 80\%) - \text{Normal RAM Usage}$$

Laptop Example:



$$\text{Max Data Limit} = (5.9 \text{ gb} \times 80\%) - 3.2 \text{ gb}$$

$$\text{Max Data Limit} = \sim 1.52 \text{ gb}$$

*R data frames actually bloats data files by ~3x

$$\text{R Data Limit} = \sim 1.52 \text{ gb} \div 3 = \sim 506.7 \text{ mb}$$

OSS R Limits: RAM

INSTANCE	CORES	RAM	DISK SIZES ¹	PRICE
M64MS	64	1,750.00 GiB	2,000 GB	\$10.34/hr
M128S	128	2,000.00 GiB	4,000 GB	\$13.34/hr

Azure's VM with largest RAM*:

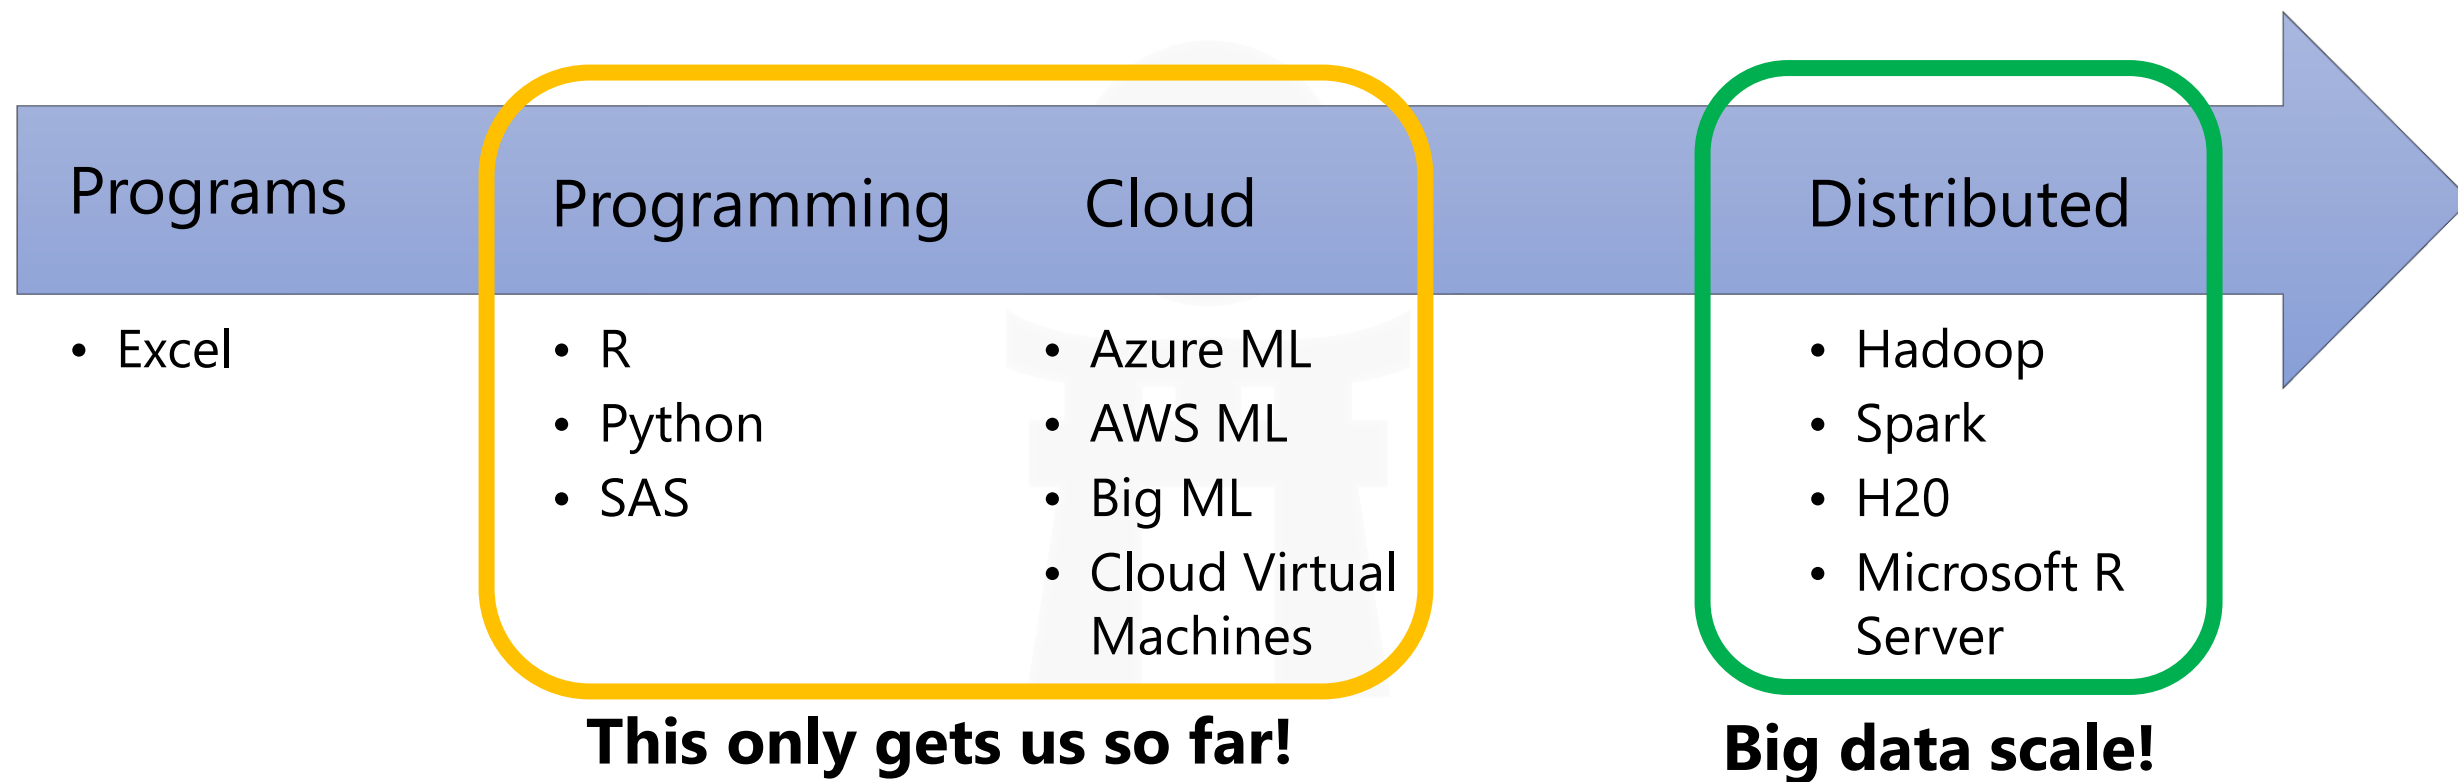
Max Data Limit = (2000gb x 80%) - 1gb

Max Data Limit = ~1600gb

R Data Limit = ~1600gb ÷ 3 = ~533.33 gb

24x7x52 Annual Cost: \$116,938.44!

Machine Learning Scaling



DISTRIBUTED COMPUTING WITH APACHE HADOOP





Turn Back The Clock, The Mainframe




- “Big Iron”
- Backbone of computing for decades.
- Still widely used.
- “Scale-up” model of shared computing.
- Core platform is cost effective, ecosystem is not (e.g., software licensing).
- The original VM host!

Distributed Computing



Cloud Computing



- Conceptually – a combination of mainframe and distributed computing.
- VM hosts are now the “Big Iron”.
- Many VMs work together to distribute workloads.
- Some workloads on dedicated HW (e.g., SAP HANA).

Scaling Computational Power




Old Scaling:

- Vertical Scaling, Scaling UP
- High performance computers




New Scaling:

- Horizontal Scaling, Scaling OUT
- Commodity hardware, distributed




What is Hadoop?

- OSS Platform for distributed computing over Internet-scale data.
- Originally built at Yahoo!
- Implementation of ideas (e.g., MapReduce) published by Google.
- The de facto standard big data platform.
- Named after a stuffed animal belonging to Doug Cutting's son.




Hadoop at Base



Distributed batch processing engine for big data.

HDFS & MapReduce



60gb of
Tweets


60gb



1 Computer

Processing: 30 hours

HDFS & MapReduce



Processing: 15 hours

HDFS & MapReduce



60 Gb of
Tweets

20Gb

20Gb

20Gb




Processing: 10 hours

3 Computers

Most Cases, Linear Scaling Of Processing Power


Number of Computers	Processing Time (hours)
1	30
2	15
3	10
4	7.5
5	6
6	5
7	4.26
8	3.75
9	3.33

If dogs were servers...






HDFS

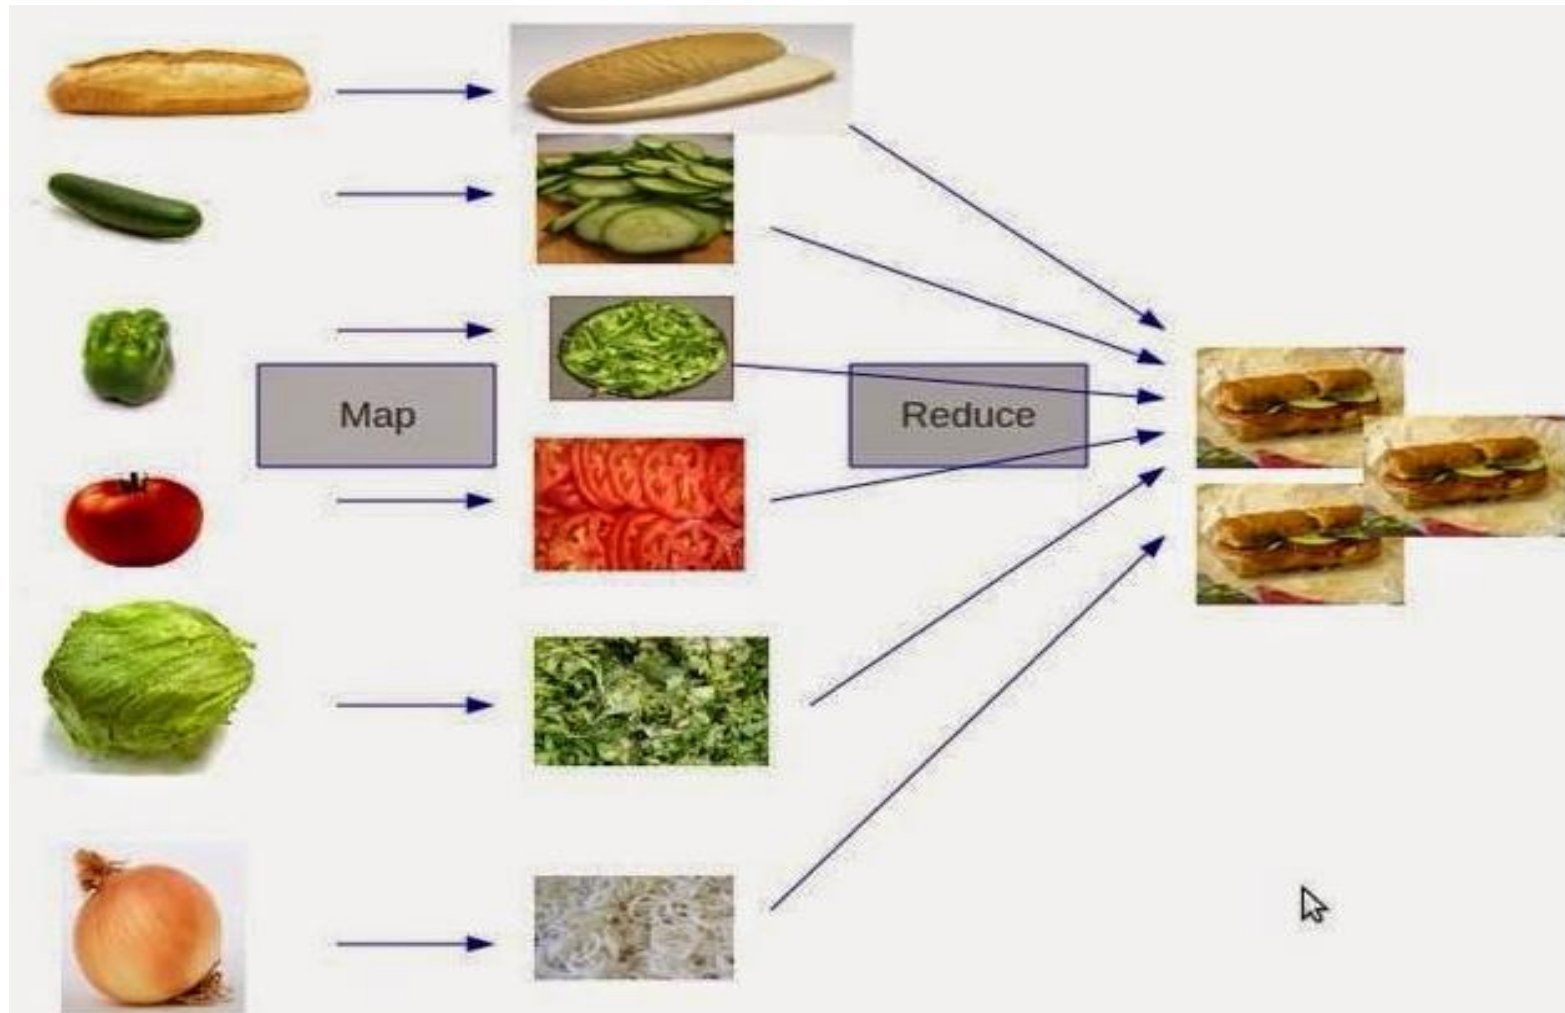
HDFS Partitioning



HDFS Redundancy



MapReduce – Sandwich Analogy



Limitations with MapReduce

- Lot of code to perform the simplest task
- Slow
- Troubleshooting multiple computers
- Good devs are scarce
- Expensive certifications

```

1  package org.apache.hadoop.examples;
2
3  import java.io.IOException;
4  import java.util.StringTokenizer;
5
6  import org.apache.hadoop.conf.Configuration;
7  import org.apache.hadoop.fs.Path;
8  import org.apache.hadoop.io.IntWritable;
9  import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Mapper;
12 import org.apache.hadoop.mapreduce.Reducer;
13 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
14 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15 import org.apache.hadoop.util.GenericOptionsParser;
16
17 public class WordCount {
18
19     public static class TokenizerMapper
20         extends Mapper<Object, Text, Text, IntWritable>{
21
22         private final static IntWritable one = new IntWritable(1);
23         private Text word = new Text();
24
25         public void map(Object key, Text value, Context context
26             ) throws IOException, InterruptedException {
27             StringTokenizer itr = new StringTokenizer(value.toString());
28             while (itr.hasMoreTokens()) {
29                 word.set(itr.nextToken());
30                 context.write(word, one);
31             }
32         }
33     }

```

DISTRIBUTED COMPUTING WITH APACHE HIVE



What is Hive?

- Abstraction built on top of MapReduce & HDFS.
- Makes Hadoop look like an RDBMS (e.g., coding in SQL).
- Developed by Facebook to democratize Hadoop.
- Applies structure to data at runtime ("schema on read").



Hive Jobs



Word Count Revisited

```

1 package org.apache.hadoop.examples;
2
3 import java.io.IOException;
4 import java.util.StringTokenizer;
5
6 import org.apache.hadoop.conf.Configuration;
7 import org.apache.hadoop.fs.Path;
8 import org.apache.hadoop.io.IntWritable;
9 import org.apache.hadoop.io.Text;
10 import org.apache.hadoop.mapreduce.Job;
11 import org.apache.hadoop.mapreduce.Mapper;
12 import org.apache.hadoop.mapreduce.Reducer;
13 import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
14 import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;
15 import org.apache.hadoop.util.GenericOptionsParser;
16
17 public class WordCount {
18     public static class TokenizerMapper
19         extends Mapper<Object, Text, Text, IntWritable>{
20
21         private final static IntWritable one = new IntWritable(1);
22         private Text word = new Text();
23
24         public void map(Object key, Text value, Context context
25             ) throws IOException, InterruptedException {
26             StringTokenizer itr = new StringTokenizer(value.toString());
27             while (itr.hasMoreTokens()) {
28                 word.set(itr.nextToken());
29                 context.write(word, one);
30             }
31         }
32     }
33 }

```

VS.

```

SELECT word,
      COUNT(*) AS word_count
FROM words
GROUP BY word;

```

Caution:

SELECT * FROM ANYTHING: This brings back everything. Everything doesn't fit on a single computer.

JOIN: Join will take hours or days to perform and eat up all cluster bandwidth for everyone else trying to use it in the queue.

ORDER BY: Sorting is very computationally expensive.


Sub Queries: A sub query essentially creates a secondary table, which will be huge in HIVE.

Interactivity: SQL in DBMS is interactive because it's almost instantaneous.

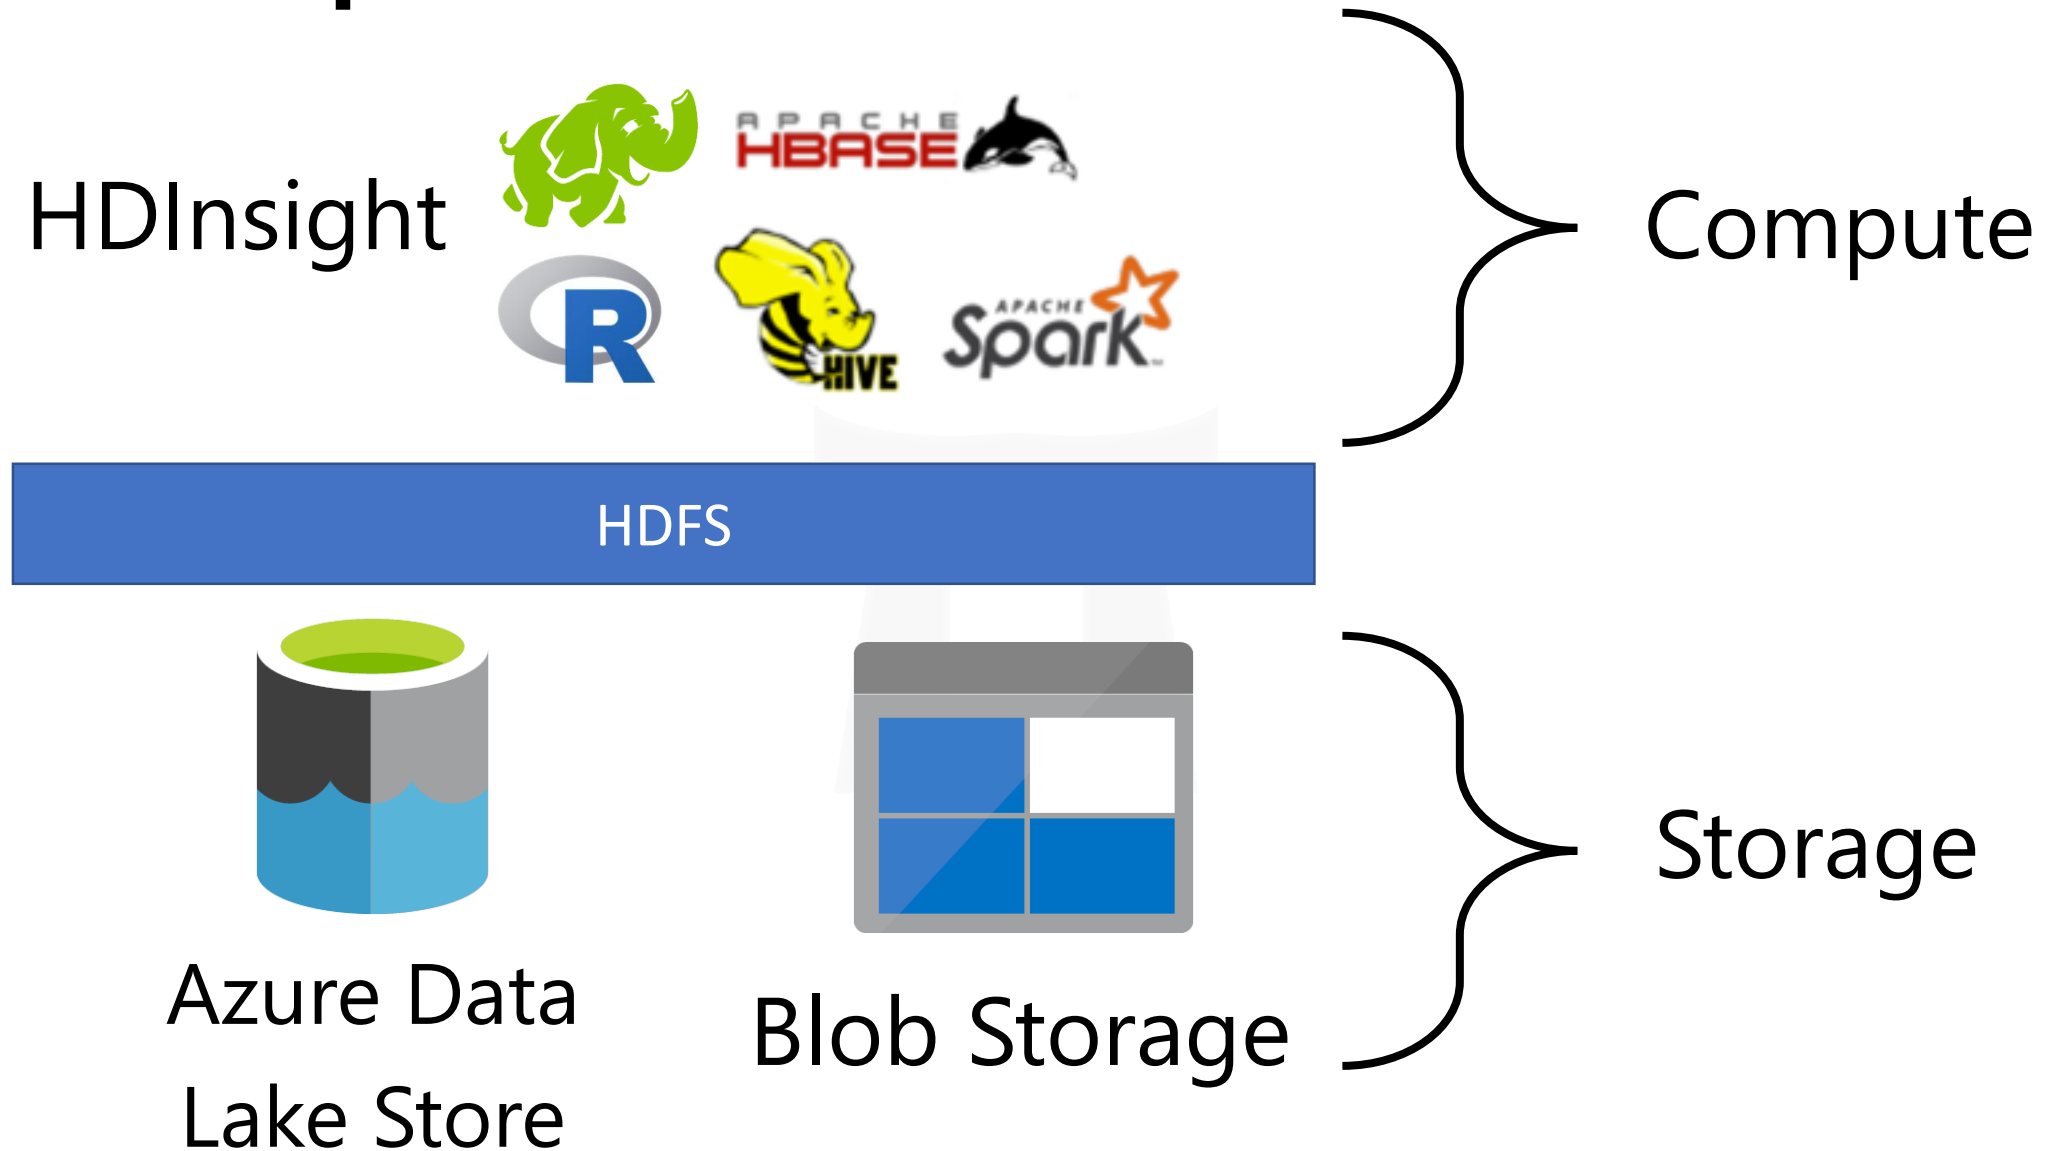
HADOOP IN THE AZURE CLOUD



Hadoop Implementations



Hadoop in Azure





MACHINE LEARNING AT SCALE - REVISITED


What is Mahout?

- Distributed Machine Learning platform.
- Built on top of MapReduce and HDFS.
- Script-based and command line interfaces.
- R-like language implementation.




Distributed Random Forest

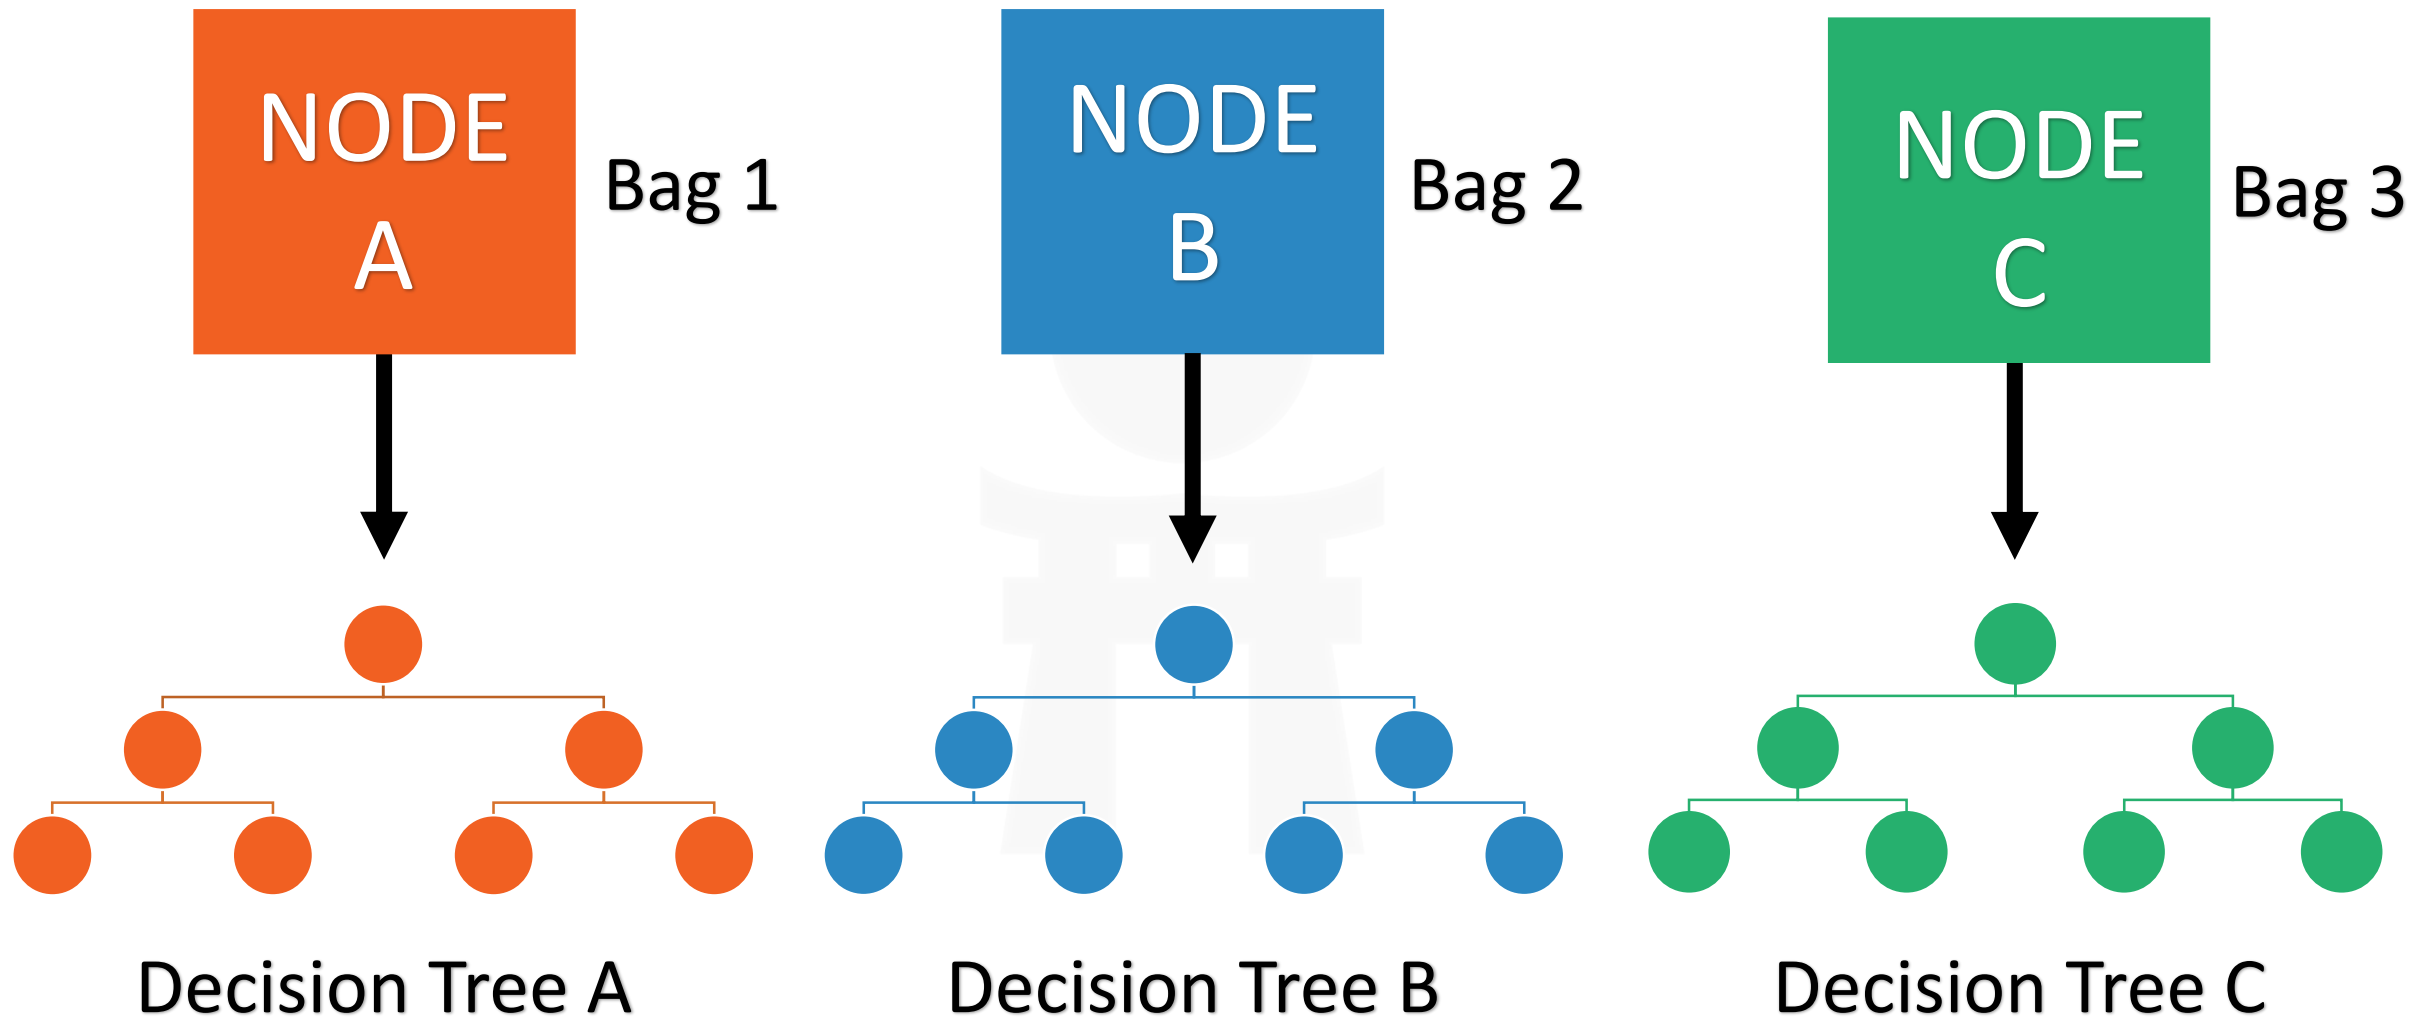
HDFS
Partitioning



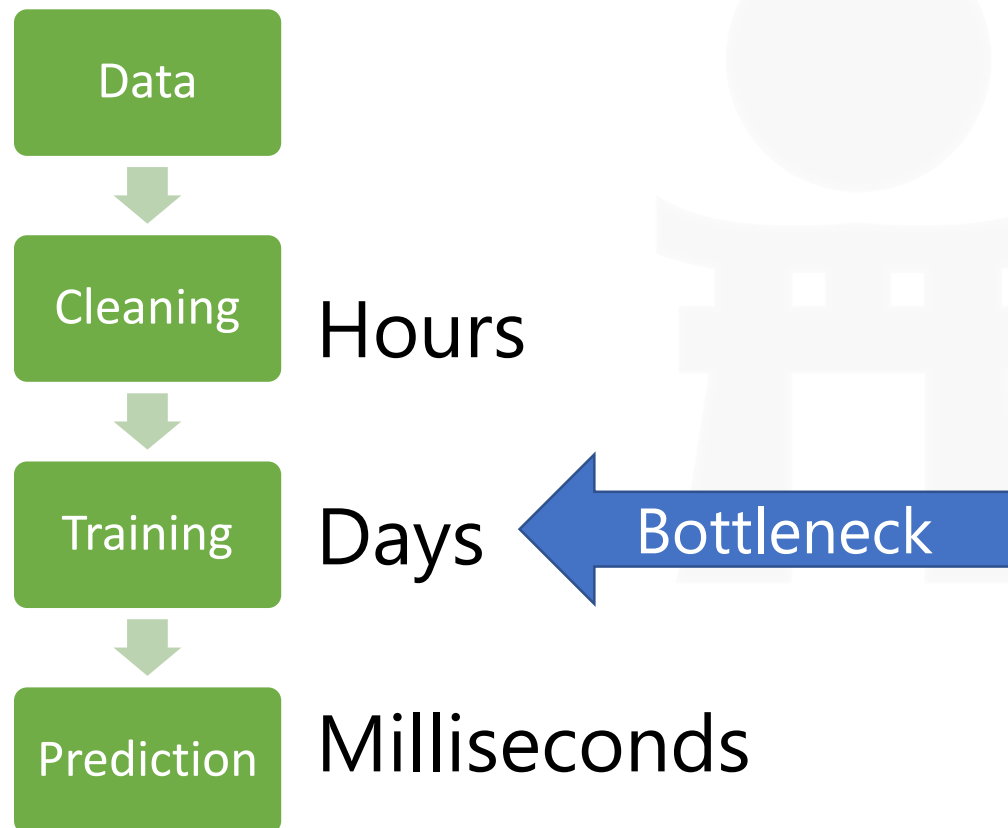
Distributed Random Forest



Distributed Random Forest



Processing Times - Machine Learning



- Large scale systems are only needed for training
- Phones can use models outputted by mahout to predict new data
- After a model is trained, save the model to any IO file type and reload it where you want



DISTRIBUTED COMPUTING V2.0 – APACHE SPARK


What is Spark?

- "A fast and general engine for large-scale data processing."
- Designed to incorporate the goodness of Hadoop and address Hadoop's shortcomings.
- Can complement Hadoop via integration with both HDFS and Hive.



Why Spark? Improved Perf!

- Up to 10x faster than Hadoop working with data from disk.*
- Up to 100x faster working with data stored in memory!*



* benchmark is without Apache Yarn

Big Data, Faster!

3x faster on 10x fewer machines!

Daytona GraySort Contest: Sort 100 TB of data!

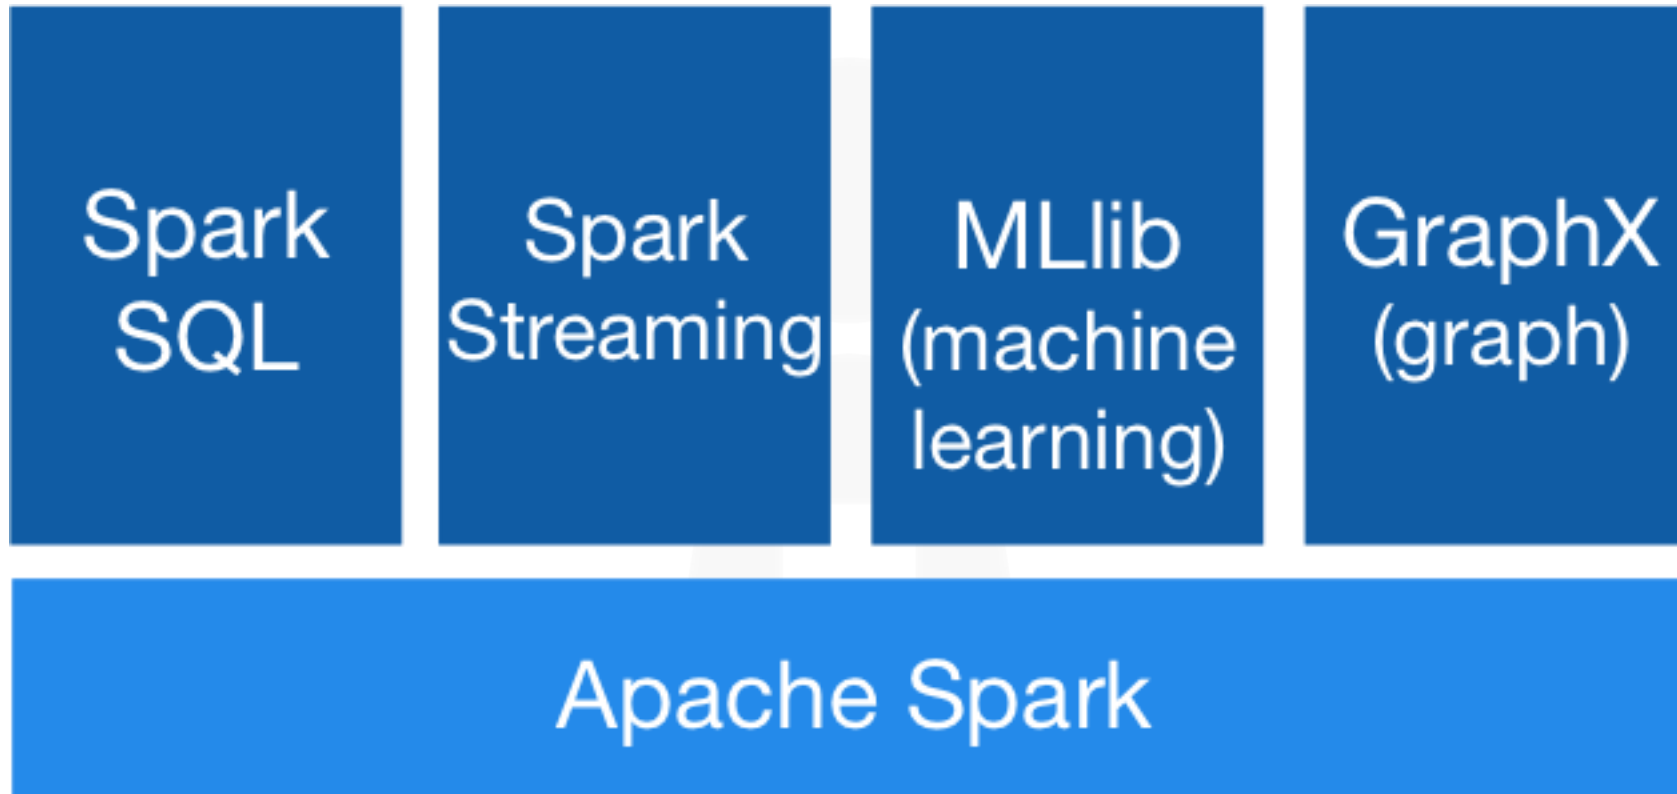
Previous World Record:

- Method: Hadoop
- Yahoo!
- 72 Minutes
- 2100 Nodes

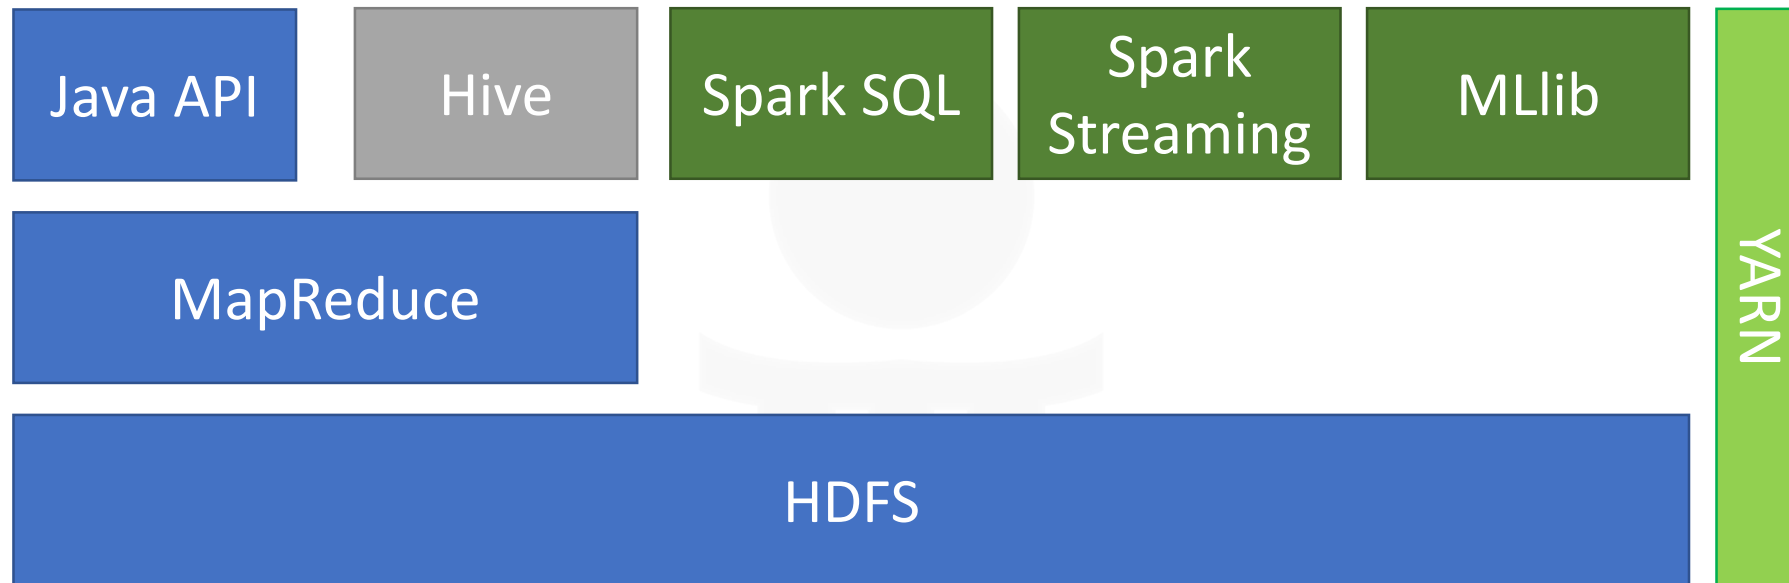
2014:

- Method: Spark
- Databricks
- 23 Minutes
- 206 Nodes

Conceptual Architecture



Spark and Hadoop




- Spark can be deployed on a Hadoop cluster and share cluster resources via YARN.
- Spark, however, does not require Hadoop!

Why is Spark Faster?

- First, Spark processing implements *lazy execution*:
 - Data operations are either *transformations* or *actions*.
 - Transformations are not executed immediately, but are stored.
 - When an action is issued, Spark evaluates all stored transformations and optimizes processing before executing.
- Second, Spark performs most processing in-memory:
 - RAM is far faster than using disk storage – even SSD drives.
 - More RAM in the cluster allows Spark to process data faster.

Technology adoption life cycle



Source: <http://carlosmartinezt.com/2010/06/technology-adoption-life-cycle/>

QUESTIONS



APPENDIX




MapReduce, via Playing Cards



Let's count the number of spades, clubs, hearts, and diamonds in a stack of cards, the way map reduce would.


- Each card represents a row of data
- Each suit & number represents an attribute of the data

Using a 2 Data Node Cluster




Mapping: Each Node's HDFS

Node 1




Node 2




Mapping: Node Sorting


Node 1



Node 2



Shuffle Sort and Data Transfer



Mapping: Node Shuffle, Data Transfer

